

Homework-4: TCP Command Server with Process Management

⚠ Important: This assignment focuses on TCP socket programming combined with process management using `fork()`. Students will create a server that forks a new child process for each command received from a client.

You must use these strict gcc compilation flags while building your application:

`-std=c17`
`-D_POSIX_C_SOURCE=200809L`
`-Wall`
`-Wextra`
`-Werror`
`-pedantic`
`-g`
`-O0`
`-pthread`

⚠ There is only one submission - RESUBMISSION IS NOT ALLOWED. You must test your implementation with the provided testing framework on the Ocelot server before the final submission.

Grader Contact Information

For any questions or concerns related to assignment grading, please reach out to the TA listed below.

Communication Guidelines:

- Ensure that all communication is polite and respectful
- **You must contact the Graders first for any grading issues**
- If the matter remains unresolved, feel free to reach out to the instructor

Name: Instructor — bbhatkal@fiu.edu (⚠ via the Canvas inbox only)

Section: COP 4610-U01 (84957)

Name: Angie Martinez — amart1070@fiu.edu (⚠ via the Canvas inbox only)

Section: COP 4610-U02 (85031)

Name: Srushti Visweswaraiah — svisw003@fiu.edu (⚠ via the Canvas inbox only)

Sections: COP 4610-UHA (85236)

Section 1: Homework Overview

To create a **TCP command server** that accepts commands from a client and executes them in separate child processes, demonstrating both network programming and process management concepts.

Section 2: Learning Objectives

- Implement TCP socket programming for client-server communication.
- Use `fork()` system call to create child processes for each command.
- Execute system commands using `system()` function.
- Understand socket communication patterns and process lifecycle management.
- Apply proper error handling in network and process contexts.

Section 3: Problem Description

You will implement a **TCP command server** that accepts connections from clients and executes commands sent by those clients. The key requirement is that the server must **fork a new child process for each command received**, not just for each client connection.

Architecture Overview

1. **Server:** Listens on port 8080 and accepts one client connection
2. **Client:** Connects to server and sends commands from an input file
3. **Process Management:** Server forks a new child process for each command
4. **Command Execution:** Child processes execute commands using `system()` and report status back

Key Features

1. **TCP Socket Communication** – Server-client communication using TCP sockets
2. **Fork Per Command** – Create a new child process for each command (not per client)
3. **Command Execution** – Use `system()` to execute commands in child processes
4. **Status Reporting** – Send command success/failure status back to client
5. **Sequential Processing** – Handle commands one by one from the client
6. **Process Management** – Proper child process creation, execution, and cleanup

Section 4: Provided Files Framework

 Refer to these files for understanding the exact requirements.

 Please do not make any modifications to the framework files provided. You are required to implement **only** the file `submission_HW4.c`. Any changes to the framework files may result in errors during evaluation or loss of credit.

Folder Structure

```

1 PROVIDED_FILES/
2   └── submission_HW4.c      # Template file (implement your solution here)
3   └── tcp_server.h          # Header file (Do not modify!)
4   └── driver.c              # Server startup framework (Do not modify!)
5   └── tcp_client.c          # Client program (Do not modify!)
6   └── Makefile                # Makefile (Do not modify!)
7   └── autograder_HW4.sh      # Testing script (Do not modify!)
8   └── batchgrader_HW4.sh      # Testing script (Do not modify!)
9   └── Sample_Executable/
10    |   └── server           # TCP server sample executable
11    |   └── client           # TCP client sample executable
12   └── Testing/               # Testcases (Do not modify!)
13     └── Testcases/
14       └── input1.txt        # Valid commands only
15       └── input2.txt        # Invalid commands only
16       └── input3.txt        # Mixed valid/invalid commands
17     └── Expected_Output/
18       └── output1.txt       # Expected output for test 1
19       └── output2.txt       # Expected output for test 2
20       └── output3.txt       # Expected output for test 3

```

Section 5: Submission Requirements

File Requirements (⚠ exact names required)

- **submission_HW4.c**: Your complete implementation
- **README.txt or README.pdf**: Student information and documentation
 - Your README ⚠ **MUST include**:

```

1 # Student Information
2 - Full Name: [Your Full Name]
3 - PID: [Your FIU Panther ID]
4 - Section: [Your Course Section]
5 - FIU Email: [Your @fiu.edu email]
6
7 # Homework: TCP Command Server
8 [Brief description of your implementation approach]
9 [Mention how you implemented the fork-per-command architecture]
10 [Describe any challenges faced and how you solved them]

```

Deliverable Folder (ZIP file) Structure - ⚠ exact structure required - no additional files/subfolders

⚠ **DO NOT submit any other files other than these required files. The batch autograder may treat any additional items in the ZIP file as invalid, which will result in a grade of zero:** submission_HW4.c, README.txt

📁 All other required framework files will be supplied by the instructor to the autograder during final grading.

- ⚠ You are required to submit your work as **a single ZIP archive file**.
- ⚠ **Filename Format:** your `Firstname_Lastname.zip` (exact format required).
- **You will be held responsible for receiving a ZERO grade if the submission guidelines are not followed.**

```

1 Harry_Potter.zip
2   └── submission_HW4.c    # Your server implementation
3   └── README.txt          # Your details and implementation approach

```

✓ Section 6: Developing and Testing Your Implementation

⚠ **DO NOT modify the following provided framework files:**

autograder_HW4.sh, batchgrader_HW4.sh, tcp_server.h, driver.c, test_client.c, Makefile, input.txt, and output.txt.

Required Functions

You must implement these functions in `submission_HW4.c`:

```
void handle_client_request(int client_socket, struct sockaddr_in
client_addr)
```

- **Input:** Client socket descriptor and client address structure
- **Process:**
 1. Accept commands from client one by one in a loop
 2. For each command received, fork a new child process
 3. Child process executes command and sends status back
 4. Parent process waits for child completion and continues
 5. Handle client disconnection gracefully
- **Output Format:**

```

1 Server: Connected to client [IP]:[PORT]
2 Server: Received command: [command]
3 Server: Child <pid> started for command: [command]
4 Server: Child <pid> completed
5 Server: Client [IP]:[PORT] disconnected

```

```
void execute_command_in_child(int client_socket, const char *command)
```

- **Input:** Client socket and command string to execute
- **Process:**
 1. Print child process information
 2. Execute command using `system()`
 3. Check return status and create appropriate message
 4. Send status message to client via socket
 5. Close socket and exit child process
- **Output Format:**

```
1 | child <pid>: Executing command: [command]
```

Command Execution Requirements

- Use `system()` function to execute commands in child processes
- **⚠️ IMPORTANT NOTE:** Your TCP server implementation must use `system()` to execute commands received from the client but should NOT capture or send the actual command output back to the client - instead, it must only send status messages indicating whether the command completed successfully or failed. When commands like `echo hello` are executed, their output (e.g., "hello") will appear on the server terminal but will NOT be sent to the client
- **Status Messages:**
 - Success (`system()` returns 0): "Command 'command' completed successfully"
 - Failure (`system()` returns non-zero): "Command 'command' failed"
- **Process Flow:**
 - Child: Execute command → Send status → Exit
 - Parent: Wait for child → Continue accepting commands

Example Input/Output

Sample Input (`input1.txt`):

```
1 | # valid commands
2 | true
3 | echo hello
4 | false
```

Expected Output Pattern:

⚠️ Note: The server terminal output is followed by the client terminal output, for all three expected outputs.

```
1 | # Server terminal output
2 | Starting TCP Command Server...
3 | Server: Listening on port 8080...
4 | Server: Waiting for client connection...
5 | Server: Connected to client 127.0.0.1:xxxxx
```

```

6 Server: Received command: true
7 Server: Child <PID> started for command: true
8 Child <PID>: Executing command: true
9 Server: Child <PID> completed
10 Server: Received command: echo hello
11 Server: Child <PID> started for command: echo hello
12 Child <PID>: Executing command: echo hello
13 hello
14 Server: Child <PID> completed
15 Server: Received command: false
16 Server: Child <PID> started for command: false
17 Child <PID>: Executing command: false
18 Server: Child <PID> completed
19 Server: Client 127.0.0.1:xxxxx disconnected
20 Server: Shutting down
21
22 # Client terminal output
23 Client: Connected to server
24 Client: Sending command: true
25 Client: Received: Command 'true' completed successfully
26 Client: Sending command: echo hello
27 Client: Received: Command 'echo hello' completed successfully
28 Client: Sending command: false
29 Client: Received: Command 'false' failed
30 Client: Disconnected from server

```

Note about Command Output:

- The line `hello` appears in the server terminal because the `echo hello` command executed successfully
- This command output is NOT sent to the client - only the status message is sent
- The client receives only: "Command 'echo hello' completed successfully"

STEP 1 - Sample Executable Testing:

 To better understand the client-server implementation requirements, you are **highly encouraged** to test the instructor-provided executables for both the server and client:

```

1 # Provide execute permissions to both server and client executables:
2 chmod 777 server client
3
4 # Run the server in one terminal
5 ./server
6
7 # Run the client in another terminal with an input test case
8 ./client Testing/Testcases/input1.txt
9
10 # Try with the other two test cases in the same way
11 ./client Testing/Testcases/input2.txt
12 ./client Testing/Testcases/input3.txt

```

STEP 2 - Manual testing:

Ensure the following files are located in the **same directory**:

- Your implementation: `submission_HW4.c`
- All the provided framework files
- The provided testcase folder: `Testing/`

```

1 # Explore the Makefile first to understand build targets and compilation settings
2 cat Makefile
3
4 # Build both server and client executables
5 make rebuild
6
7 # Run server in one terminal
8 ./server
9
10 # Run the client in another terminal
11 # Test with:
12 #   - Valid commands:    input1.txt
13 #   - Invalid commands: input2.txt
14 #   - Mixed commands:   input3.txt
15 ./client Testing/Testcases/input1.txt
16 ./client Testing/Testcases/input2.txt
17 ./client Testing/Testcases/input3.txt
18
19 # Prepare the student output files by copy-pasting the outputs from both
20 # the server terminal and the client terminal into the following files:
21 student_output1.txt
22 student_output2.txt
23 student_output3.txt
24
25 ⚠️ Normalizing the process IDs and PORT numbers in student_output.txt.
26 # (The autograder replaces all actual PIDs with <PID> before comparison.)
27 # For example, your output:
28 #   - Process IDs: <1234> becomes <PID>
29 # will match the expected pattern:
30 #   - Port numbers: 127.0.0.1:54321 becomes 127.0.0.1:xxxxx
31
32 # Compare your output with the expected output:
33 diff student_output1.txt Testing/Expected_Output/output1.txt
34 diff student_output2.txt Testing/Expected_Output/output2.txt
35 diff student_output3.txt Testing/Expected_Output/output3.txt

```

STEP 3 - Autograder testing:

⚠️ Important: ⚠️ The autograder may take longer to test your application due to client-server interactions. Please be patient!

⚠️ The autograder may take longer to test your application due to client-server interactions. Please be patient!

Ensure the following files are located in the **same directory**, and then run the **Autograder**:

- **Your implementation:** `submission_HW4.c`
- **All the provided framework files**
- **The provided testcase folder:** `Testing/`

```
1 # Run the autograder  
2 ./autograder_HW4.sh
```

- **Check your grade** and fix any issues
- **Repeat until** you achieve the desired score

STEP 4 - Batch Autograder testing:

⚠ Important: ⚠ The batch autograder may take even longer than the autograder to test your application due to additional testing steps. Please be patient!

The batch autograder, **used by the instructor**, processes all student submissions at once. It utilizes the provided framework files, extracts the required files from your submitted `.zip` file, and performs grading accordingly.

⚠ Backup Files Before Running Batch Grader - The script deletes existing files to rebuild the environment from scratch.

Ensure the following files are located in the **same directory**, and then run the **Batch Autograder**:

- **All the provided framework files**
- **The provided testcase folder:** `Testing/`
- **Your final submission ZIP file** consisting the required files for the submission: example, `Harry_Potter.zip`

```
1 # Run the batch-autograder  
2 ./batchgrader_HW4.sh
```

Final Testing requirements:

- ⚠ **MUST test on ocelot server:** `ocelot-bbhatkal.aul.fiu.edu`
- ⚠ **Test thoroughly before submission** - no excuses accepted
- ⚠ **"Works on my computer"** is NOT accepted as an excuse
- **If you pass all test cases on the server, you will likely pass instructor's final grading**
- **What you see is what you get** - autograder results predict your final grade

Section 7: Grading Criteria

Autograder Testing

- ⚠ Your implementation `submission_HW4.c` will be tested against the provided test case as well as additional instructor test cases using the autograder.
- ⚠ **Exact output matching required** - any deviation results in point deduction
- ⚠ **Program must compile** without errors or warnings

Test Case Distribution:

- Test 1 (30 points):** Valid commands only - tests successful process execution and status reporting
- Test 2 (30 points):** Invalid commands only - tests error handling for non-existent commands
- Test 3 (40 points):** Mixed valid/invalid commands - tests comprehensive functionality including:
 - Proper fork-per-command behavior
 - Correct client-server communication
 - Process management and cleanup
 - Recovery after command failures

Penalties

- ⚠ **Missing README: -10 points**
- ⚠ **Missing `submission_HW4.c` ZERO grade** (autograder compilation failure)
- ⚠ **Incorrect ZIP filename: ZERO grade** (autograder compilation failure)
- ⚠ **Wrong source filename: ZERO grade** (autograder compilation failure)
- ⚠  **Resubmission: NOT ALLOWED**

Section 8: Technical Specifications

Compilation Requirements

- Build System:** A complete Makefile is provided as part of the framework
- Compiler Standards:** The Makefile uses strict C17 compliance with POSIX.1-2008 extensions
- Important:** Explore the provided Makefile to understand the compilation flags and build targets
- Cross-platform:** The build configuration ensures compatibility across WSL, Ubuntu, and macOS
- Implementation Checks:** The Makefile includes automated checks to verify your implementation contains required functions

Socket Programming Requirements

- Server listens on **port 8080**
- Accept **one client connection** at a time
- Use **TCP sockets** (SOCK_STREAM)
- Handle client disconnection gracefully

Process Management Requirements

- **Fork for each command** received (not per client)
- Child processes use **system()** to execute commands
- Parent process **waits** for child completion using **waitpid()**
- Proper **zombie process** reaping

Communication Protocol

- Client sends commands as strings
- Server receives commands and forks for execution
- Child processes execute commands using system() - command output appears on server terminal
- Child sends **only status message** back to client (not command output)
- Client displays received status messages

Error Handling

- Check **fork()** return value for failure
- Handle **socket errors** appropriately
- Manage **command execution failures**
- Implement **timeout protection** in autograder

Section 9: Academic Integrity

-  **This is an individual assignment**
-  **You may discuss concepts** but not share code
-  **All submitted code** must be your original work
-  **You are encouraged to adapt from instructor provided socket examples** but must understand and modify them appropriately
-  **Plagiarism** is a serious offense and will result in penalties - **ZERO grade** (academic integrity violation).

 **Good luck with your TCP command server implementation!**