

Homework-6: Round Robin (RR) CPU Scheduler

⚠ Important: This homework focuses on implementing the Round-Robin (RR) CPU scheduling algorithm. Students will work with process scheduling concepts including arrival times, burst times, time quantum, waiting times, and turnaround times.

💡 All the necessary concepts required to complete this homework are covered in the [CPU Scheduling lecture materials and textbook chapter on scheduling algorithms](#).

You must use these strict gcc compilation flags while building your application:

**-std=c17
-D_POSIX_C_SOURCE=200809L
-Wall
-Wextra
-Werror
-g
-O0**

⚠ There is only one submission - RESUBMISSION IS NOT ALLOWED. You must test your implementation with the provided testing framework on the Ocelot server before the final submission.

Grader Contact Information

For any questions or concerns related to assignment grading, please reach out to the TA listed below.

Communication Guidelines:

- Ensure that all communication is polite and respectful
- **You must contact the Graders first for any grading issues**
- If the matter remains unresolved, feel free to reach out to the instructor

Name: Instructor – bbhatkal@fiu.edu (⚠ via the Canvas inbox only)

Section: COP 4610-U01 (84957)

Name: Angie Martinez – amart1070@fiu.edu (⚠ via the Canvas inbox only)

Section: COP 4610-U02 (85031)

Name: Srushti Visweswaraiah – svisw003@fiu.edu(⚠ via the Canvas inbox only)

Sections: COP 4610-UHA (85236)

Section 1: Homework Overview

To implement the Round-Robin (RR) CPU scheduling algorithm that schedules processes using a fixed time quantum in a circular queue fashion, demonstrating preemptive CPU scheduling concepts, context switching, and metrics calculation.

Section 2: Learning Objectives

- Understand and implement the Round-Robin (RR) CPU scheduling algorithm.
- Work with process control blocks containing scheduling information.
- Implement preemptive scheduling with time quantum.
- Manage a ready queue using circular queue principles.
- Calculate turnaround time and waiting time for processes.
- Implement proper tie-breaking rules when processes have the same arrival time.
- Handle CPU idle time when no processes are ready to execute.
- Track context switches in preemptive scheduling.

Section 3: Problem Description

You will implement the **Round-Robin (RR) scheduling algorithm** that schedules processes using a fixed time quantum in a circular manner. RR is a preemptive scheduling algorithm where each process gets a small unit of CPU time (time quantum), after which it is preempted and added to the end of the ready queue.

Algorithm Overview

1. **Ready Queue Management:** Maintain a ready queue of processes waiting for CPU time
2. **Time Quantum:** Each process executes for at most one time quantum before being preempted
3. **Preemption:** If a process doesn't complete within its time quantum, it's moved to the back of the ready queue
4. **Process Arrival Handling:** New arrivals are added to the ready queue in order of arrival time
5. **Tie-Breaking:** If multiple processes arrive at the same time, use Process ID (lower PID first)
6. **CPU Idle Time:** If no process is ready, CPU remains idle until the next process arrives
7. **Metrics Calculation:** Calculate waiting time and turnaround time for each process

Key Concepts

Time Quantum (TQ): Fixed time slice allocated to each process (provided as input)

Turnaround Time (TAT): Total time from arrival to completion

- Formula: `TAT = Completion Time - Arrival Time`

Waiting Time (WT): Time a process spends waiting in the ready queue

- Formula: `WT = Turnaround Time - CPU Burst Time`

Average Turnaround Time (ATT): Mean of all processes' turnaround times

- Formula: `ATT = Sum of all TAT / Number of Processes`

Average Waiting Time (AWT): Mean of all processes' waiting times

- Formula: `AWT = Sum of all WT / Number of Processes`

⚠ Note: All these time metrics are automatically computed by the framework. Students do not need to implement this functionality.

Round-Robin Execution Rules

1. Initial Queue Setup: At time 0, all processes with arrival time 0 enter the ready queue (sorted by PID if tied)

2. Process Execution:

- The first process in the ready queue executes for $\min(\text{time_quantum}, \text{remaining_time})$
- If $\text{remaining_time} \leq \text{time_quantum}$, the process completes
- If $\text{remaining_time} > \text{time_quantum}$, the process is preempted

3. New Arrivals During Execution:

- Check for processes that arrive during the current execution window
- Add newly arrived processes to the ready queue (in order of arrival, then PID)

4. Queue Reordering After Execution:

- If the current process was preempted, add it to the end of the ready queue
- Move to the next process in the ready queue

5. CPU Idle Handling: If the ready queue is empty, advance time to the next process arrival

Tie-Breaking Rules

RR Tie-Breaking: When multiple processes have the same arrival time:

1. **Primary:** Arrival Time (ascending order - earlier first)

2. **Secondary:** Process ID (ascending order - lower PID first)

Example: If P3 and P1 both arrive at time 0, P1 enters the ready queue first (lower PID)

Section 4: Provided Files Framework

 **Refer to these files** for understanding the exact requirements.

⚠ Please do not make any modifications to the framework files provided. You are required to implement **only** the file `submission_HW6.c`. Any changes to the framework files may result in errors during evaluation or loss of credit.

Folder Structure

```

1 PROVIDED_FILES/
2   |- skeleton_submission_HW6.c      # Implementation template file (rename to submission_HW6.c)
3   |- rr_scheduler.h                # Header file (Do not modify!)
4   |- driver.c                     # Main driver program (Do not modify!)
5   |- Makefile                      # Build system (Do not modify!)
6   |- autograder_HW6.sh            # Testing script (Do not modify!)

```

```

7  └── batchgrader_HW6.sh          # Batch testing script (Do not modify!)
8  └── Sample_Executable/
9    └── rr                         # Sample executable for testing
10   └── Testing/
11     └── Testcases/
12       ├── input1.txt             # Test case 1
13       ├── input2.txt             # Test case 2
14       ├── input3.txt             # Test case 3
15       ├── input4.txt             # Test case 4
16       ├── input5.txt             # Test case 5
17       ├── input6.txt             # Test case 6
18       ├── input7.txt             # Test case 7
19       └── input8.txt             # Test case 8
20
21   └── Expected_Output/
22     ├── output1.txt            # Expected output for test 1
23     ├── output2.txt            # Expected output for test 2
24     ├── output3.txt            # Expected output for test 3
25     ├── output4.txt            # Expected output for test 4
26     ├── output5.txt            # Expected output for test 5
27     ├── output6.txt            # Expected output for test 6
28     ├── output7.txt            # Expected output for test 7
29     └── output8.txt            # Expected output for test 8

```

Section 5: Submission Requirements

File Requirements (⚠ exact names required)

- **submission_HW6.c**: Your complete Round-Robin implementation
- **README.txt or README.pdf or README.md**: Student information and documentation
 - Your README ⚠ **MUST include**:

```

1  # Student Information
2  - Full Name: [Your Full Name]
3  - PID: [Your FIU Panther ID]
4  - Section: [Your Course Section]
5  - FIU Email: [Your @fiu.edu email]
6
7  # Homework: Round-Robin CPU Scheduler
8  [Brief description of your implementation approach]
9  [Explain how you implemented the Round-Robin scheduling algorithm]
10 [Describe how you managed the ready queue and time quantum]
11 [Explain how you handled process arrivals during execution]
12 [Describe how you handled tie-breaking when processes have same arrival time]
13 [Mention any challenges faced and how you solved them]

```

Deliverable Folder (ZIP file) Structure - ⚠ exact structure required - no additional files/subfolders

⚠ **DO NOT submit any other files other than these required files. The batch autograder may treat any additional items in the ZIP file as invalid, which will result in a grade of zero:** submission_HW6.c, README.txt (or README.pdf or README.md)

└ All other required framework files will be supplied by the instructor to the autograder during final grading.

- ⚠ You are required to submit your work as **a single ZIP archive file**.
- ⚠ **Filename Format:** your `Firstname_Lastname.zip` (exact format required).
- **You will be held responsible for receiving a ZERO grade if the submission guidelines are not followed.**

```

1 Harry_Potter.zip
2 └─ submission_HW6.c # Your Round-Robin scheduler implementation
3 └─ README.txt       # Your details and implementation approach

```

✓ Section 6: Developing and Testing Your Implementation

⚠ **DO NOT modify the following provided framework files:**
`autograder_HW6.sh`, `batchgrader_HW6.sh`, `rr_scheduler.h`, `driver.c`, `Makefile`, and `Testing/ folder`.

Required Function

You must implement this function in `submission_HW6.c`:

```
void round_robin_scheduling(SchedulerContext *ctx, int time_quantum)
```

- **Input:**
 - Pointer to SchedulerContext structure containing process information
 - Integer `time_quantum` representing the time slice for each process
- **Process:**
 1. Reset process states using the 💡 provided `reset_process_states(ctx)`
 2. Initialize a ready queue to manage processes
 3. Sort processes by arrival time (use Process ID for tie-breaking)
 4. Simulate Round-Robin scheduling:
 - Track current time
 - Manage ready queue (add new arrivals, handle preemptions)
 - Execute processes for at most `time_quantum` units
 - Handle CPU idle time when needed
 - Calculate waiting time for each process

5. Display results using **provided** `display_results(ctx, "RR")`

- **Important Notes:**

- Do NOT modify the global `ctx->processes[]` array order
- Create a local copy/queue for scheduling
- Update waiting times in the original global array
- Properly handle process arrivals that occur during execution
- Track remaining time for each process

Input File Format

Each test case file follows this format:

```

1 =====
2 TEST CASE 1: All Arrive at Time 0 (TQ=4)
3 Purpose: Basic Round Robin with multiple processes requiring multiple rounds
4 =====
5 Process    Burst Time    Priority    Arrival Time
6 =====
7 P1          6             0           0
8 P2          5             0           0
9 P3          8             0           0
10 P4         7             0           0

```

Format Explanation:

- **Lines 1-4:** Test case header and description (informational only)
- **Line 5:** Column headers (Process, Burst Time, Priority, Arrival Time)
- **Line 6:** Separator line (equal signs)
- **Lines 7+:** Process data rows
 - **Column 1:** Process ID (format: P1, P2, P3, etc.)
 - **Column 2:** Burst Time (CPU time required)
 - **Column 3:** Priority (not used in RR, but present in data)
 - **Column 4:** Arrival Time (when process enters ready queue)

Note:

- Your program will receive the time quantum (TQ=4) through the driver
- Priority field is ignored in Round-Robin scheduling but will be used in future assignments
- The driver automatically skips header lines and reads only the process data

Output Format

Your implementation must produce output in this exact format:

```

1 Round-Robin (RR):
2 PID      Turnaround_Time      Waiting_Time
3 1          3                  0
4 2          10                 5
5 3          9                  7
6 Average Turnaround Time: 7.33
7 Average Waiting Time: 4.00

```

Output Format Explanation:

- **Line 1:** Algorithm name ("Round-Robin")
- **Lines 2-7:** Process results (one per process, in original input order)
 - **Column 1:** Process ID (PID)
 - **Column 2:** Turnaround Time for that process
 - **Column 3:** Waiting Time for that process
 - **Line 6:** Average Turnaround Time
 - **Line 7:** Average Waiting Time

STEP 1 - Sample Executable Testing:

 To better understand the Round-Robin implementation requirements, you are **highly encouraged** to test the instructor-provided executable:

```

1 # Navigate to Sample_Executable folder
2 cd Sample_Executable
3
4 # Provide execute permissions
5 chmod 777 rr
6
7 # Run with test cases (redirect input from test files)
8 # Repeat the process for all 8 test cases by changing the input file name accordingly.
9 ./rr < ../Testing/Testcases/input1.txt
10
11 # Compare output with expected results
12 ./rr < ../Testing/Testcases/input1.txt > my_output1.txt
13 diff my_output1.txt ../Testing/Expected_Output/output1.txt

```

STEP 2 - Manual testing:

Ensure the following files are located in the **same directory**:

- **Your implementation:** `submission_HW6.c`
- **All the provided framework files**
- **The provided testcase folder:** `Testing/`

```

1 # Explore the Makefile first to understand build targets and compilation settings
2 cat Makefile

```

```

3
4 # Build the executable
5 make clean
6 make
7
8 # Test with all eight test cases
9
10 ./rr < ../Testing/Testcases/input1.txt
11
12
13 # Save your output for comparison
14 # Repeat the process for all 8 test cases by changing the input file name
15 # and student_output file name accordingly.
16 ./rr < Testing/Testcases/input1.txt > student_output1.txt
17
18 # Compare your output with expected output
19 # Repeat the process for all 8 test cases by changing the student_output file name
20 # and expected output file name accordingly.
21 diff student_output1.txt Testing/Expected_Output/output1.txt
22
23 # If diff shows no output, your implementation is correct!

```

STEP 3 - Autograder testing:

Ensure the following files are located in the **same directory**, and then run the **Autograder**:

- **Your implementation:** `submission_HW6.c`
- **All the provided framework files**
- **The provided testcase folder:** `Testing/`

```

1 # Make autograder executable
2 chmod +x autograder_HW6.sh
3
4 # Run the autograder
5 ./autograder_HW6.sh

```

- **Check your grade** and fix any issues
- **Repeat until** you achieve the desired score

STEP 4 - Batch Autograder testing:

⚠ Important: The batch autograder may take a bit of time. Please be patient!

The batch autograder, **used by the instructor**, processes all student submissions at once. It utilizes the provided framework files, extracts the required files from your submitted `.zip` file, and performs grading accordingly.

⚠ Backup Files Before Running Batch Grader - The script deletes existing files to rebuild the environment from scratch.

Ensure the following files are located in the **same directory**, and then run the **Batch Autograder**:

- All the provided framework files
- The provided testcase folder: `Testing/`
- Your final submission `ZIP file` consisting the required files for the submission: example, `Harry_Potter.zip`

```

1 # Make batch autograder executable
2 chmod +x batchgrader_HW6.sh
3
4 # Run the batch-autograder
5 ./batchgrader_HW6.sh

```

Final Testing requirements:

- **⚠ MUST test on ocelot server:** `ocelot-bbhatkal.aul.fiu.edu`
- **⚠ Test thoroughly before submission** - no excuses accepted
- **⚠ "Works on my computer" is NOT accepted** as an excuse
- **✓ If you pass all test cases on the server, you will likely pass instructor's final grading**
- **✓ What you see is what you get** - autograder results predict your final grade

Section 7: Grading Criteria

Autograder Testing

- **⚠ Your implementation `submission_HW6.c`** will be tested against the provided test cases as well as additional instructor test cases using the autograder.
- **⚠ Exact output matching required** - any deviation results in point deduction
- **⚠ Program must compile** without errors or warnings

Test Case Distribution:

- **Test 1 (15.5 points):** Basic Round Robin with multiple processes requiring multiple rounds
- **Test 2 (15.5 points):** Processes arriving at different times with large gaps causing CPU idle periods
- **Test 3 (15.5 points):** New process arrives exactly when current process's time quantum expires
- **Test 4 (15.5 points):** Test CPU idle time handling with large gaps between process arrivals
- **Test 5 (15.5 points):** Mix of processes with burst times both less than and greater than time quantum
- **Test 6 (15.55 points):** Edge case with only one process requiring multiple time quantum
- **Test 7 (15.5 points):** Multiple processes arrive while first process is executing within its time quantum
- **Test 8 (15.5 points):** All processes have burst time less than time quantum, testing immediate completion without preemption

🚫 Penalties

- ⚠️ Missing README: **-10 points**
- ⚠️ Missing submission_HW6.c : **ZERO grade** (autograder compilation failure)
- ⚠️ Incorrect ZIP filename: **ZERO grade** (autograder extraction failure)
- ⚠️ Wrong source filename: **ZERO grade** (autograder compilation failure)
- 🔒 Resubmission: **NOT ALLOWED**

✓ Section 8: Technical Specifications

Data Structures

You will work with these structures defined in `rr_scheduler.h`:

```

1  typedef struct {
2      int pid;           // Process ID
3      int priority;     // Priority (not used in RR)
4      int burst_time;    // CPU burst time
5      int arrival_time;  // Arrival time
6      int remaining_time; // Remaining time (YOU must update this)
7      int waiting_time;   // Waiting time (YOU calculate this)
8      int turnaround_time; // Turnaround time (framework calculates)
9      int completion_time; // Completion time
10     bool is_completed;   // Completion flag
11 } Process;
12
13 typedef struct {
14     Process processes[MAX_PROCESSES]; // Array of processes
15     int num_processes;             // Number of processes
16 } SchedulerContext;

```

✓ Section 9: Academic Integrity

- 👤 This is an individual assignment
- 💬 You may discuss concepts but not share code
- 📝 All submitted code must be your original work
- 📘 You are encouraged to refer to textbook and lecture materials but must write your own implementation
- ⚠️ Plagiarism is a serious offense and will result in penalties - **ZERO grade** (academic integrity violation).

🚀 Good luck with your Round-Robin CPU Scheduler implementation!